

Sequence Learning with Connectionist Temporal Classification

Alex Graves, 2006

Rakesh Rao
Stanford University

June 14, 2015

References

- Graves, Alex. Supervised sequence labelling with recurrent neural networks. Vol. 385. Heidelberg: Springer, 2012.
- Graves, Alex, et al. "A novel connectionist system for unconstrained handwriting recognition." Pattern Analysis and Machine Intelligence, IEEE Transactions on 31.5 (2009): 855-868.
- Graves, Alex, et al. "Connectionist temporal classification: labelling unsegmented sequence data with recurrent neural networks." Proceedings of the 23rd international conference on Machine learning. ACM, 2006.

Motivation

Learn a sequence of labels from an input stream


- Input and output of variable lengths
- Location unknown/undefined



Figure : Statistics

The Model

- Input is a sequence in \mathbb{R}^n (here $n = 20$, the image height)
- Output is a sequence in \mathbb{R}^k (here $k = 26$, the number of classes)
- k -vector sums to unity



---sssssss-----tttttt--

Recurrent Neural Network

- Say, in previous slide, $20 = 2$ and $26 = 2$, i.e. image height is two pixels, and there are only two classes. Then we can ...




Figure : Parallel lines share weights

Recurrent Neural Network

- Now each input **node** is n vector and output **node** k vector
- Each arrow is now a matrix multiplication
- Parallel arrows have same weight matrix




Figure : RNN for a sequence with seven timesteps

A Good Output for CAB

- Alphabet = $\{A, B, C, D\}$, $k = 4 + 1$
- Add a blank or null class the network can fall back to. It reduces memory burden on the network and allows repeated labels




Figure : RNN outputs for a sequence with seven timesteps. `_C.A.B_` is the most likely labeling given the output.

Looking for the intractable CAT

- Out of the t time steps, the letters we want can 'stand out' anytime.
- Each such sequence is called a *path*
e.g:- For $t = 7$ and output = cat, we can have $_ca_t_$, $_c_a_t_$, $_c_a_t_$, $ca_t_$, etc. are all *good* paths
- Intractable number of paths for large t




Figure : All six excitations give us a strong **cat** ($t = 100$)

Given all the above

- How do we train the network?
- How do we tell the cat?

Notation

- Alphabet $A' = A \cup \{\text{blank}\}$
- y_k^t - activation of k^{th} class at time t (interpreted as probability)
- A'^T - set of length T sequences over A'
- π one such *path* in A'^T
e.g. $\pi = _ca_tt_$ for $T = 7$.
- According to the model,

$$p(\pi|\mathbf{x}) = \prod_{t=1}^T y_{\pi_t}^t$$

- $\mathcal{F} : A'^T \rightarrow A^{\leq T}$ mapping from path to labeling.
e.g:- $\mathcal{F}(_ca_t_) = \mathcal{F}(ccaa_t_) = \dots = cat$
- Probability of a/correct labelling:

$$p(\mathbf{l}|\mathbf{x}) = \sum_{\pi \in \mathcal{F}^{-1}(\mathbf{l})} p(\pi|\mathbf{x})$$

All the cat's paths




Figure : Black circles are blanks, white are labels. Arrows are allowed transitions. One traversal from left to right is a *path* corresponding to the labelling cat or equivalently `_c_a_t_` (wlog)

Forward probabilities

U is length of \mathbf{l} (i.e. height of the picture in previous slide)

$$\alpha(t, u) = \sum_{\pi \in V(t, u)} \prod_{i=1}^t y_{\pi_i}^i$$

where $V(t, u)$ is the set of all paths going through label u at time t .
 (t, u) a circle in picture.


$$p(\mathbf{l}|\mathbf{x}) = \alpha(T, U) + \alpha(T, U - 1)$$

$$\alpha(1, 1) = y_{\text{blank}}^1$$

$$\alpha(1, 2) = y_{l_1}^1$$

$$\alpha(1, 3) = \alpha(1, 4) = \dots = \alpha(1, U) = 0$$

CTC will bell the CAT



Forward recursion: Just add the paths entering a circle and multiply the sum by the activation of that circle

$$\alpha(t+1, u) = \{\alpha(t, u) + \alpha(t, u-1) + \mathbb{1}(l_u \neq \text{blank}) \alpha(t, u-2)\} y_{l_u}^{t+1}$$

Summary

- Apply RNN on input

$$\mathbf{y} = \text{RNN}(\mathbf{x}; \Theta)$$

- Find Forward probabilities

$$\alpha(t+1, u) = y_{l_u}^{t+1} [\alpha(t, u) + \alpha(t, u-1) + \mathbb{1}(l_u \neq \text{blank}) \alpha(t, u-2)]$$

- Find probability of the desired labelling


$$p(\mathbf{l}|\mathbf{x}) = \alpha(T, U) + \alpha(T, U-1)$$

- Calculate Negative log-likelihood loss over the entire dataset S

$$\mathcal{L}(S) = - \sum_{(\mathbf{x}, \mathbf{l}) \in S} \ln p(\mathbf{l}|\mathbf{x})$$

- Back-propagate Symbolic-differentiate \mathcal{L} and gradient descend in the weight space for the argmin $\Theta \equiv \{\mathbf{W}_{ih}, \mathbf{W}_{hh}, \mathbf{W}_{ho}\}$

Really? Like, did it ever even, like, actually work, like, at all?




```

Input Dim: 8
Num Classes: 6
Num Samples: 1000

Preparing the Data
Building the Network
Training the Network
Epoch : 0
## TRAIN cost: 38.611
Shown : 5 2 2
Seen : 4 3 4 4 3 2 3 4 3 4 2 4 3
Image Shown:
0 |
1 |
2 |
3 |
4 |
5 |
6 |
7 |
  5 2 2
Firings:
0 |
1 |
2 |
3 |
4 |
5 |
6 |

```







Shown : 0 3 2 0 4

Seen : 2 4 4

Image Shown:

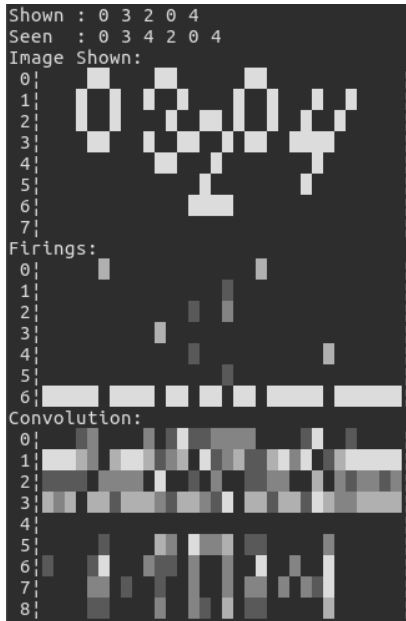


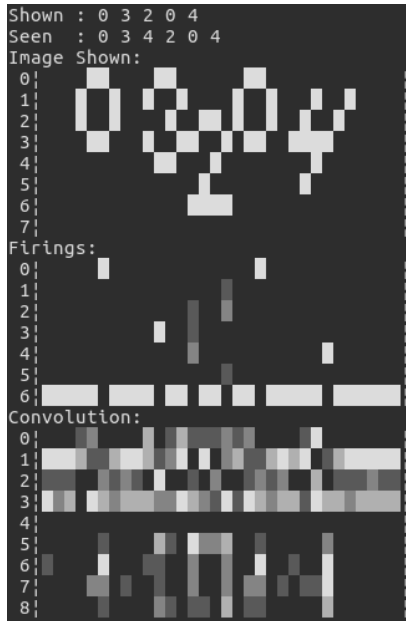
Firings:

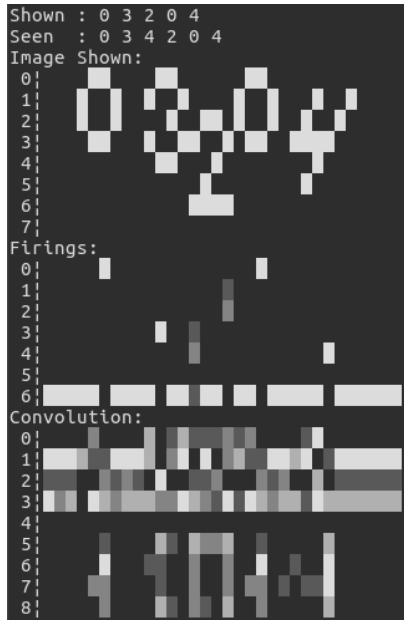


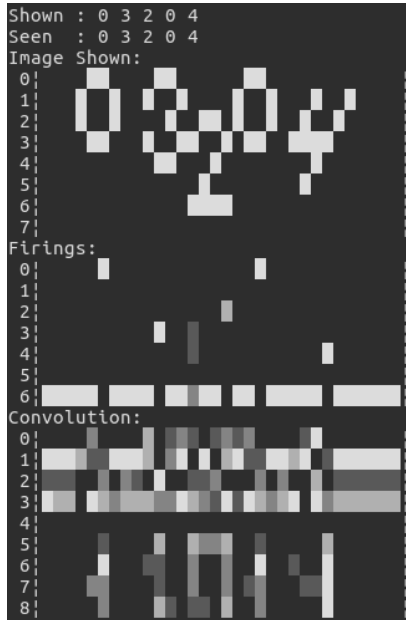
Convolution:

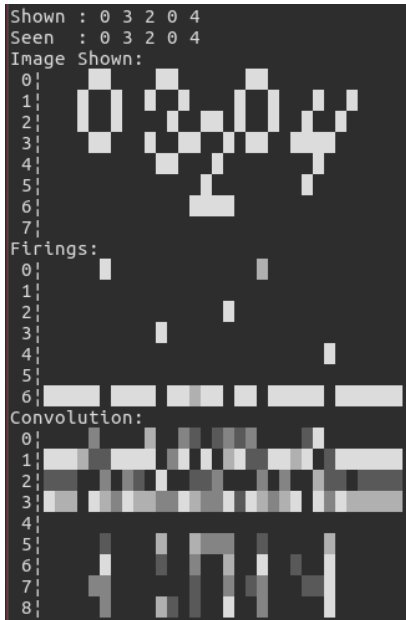


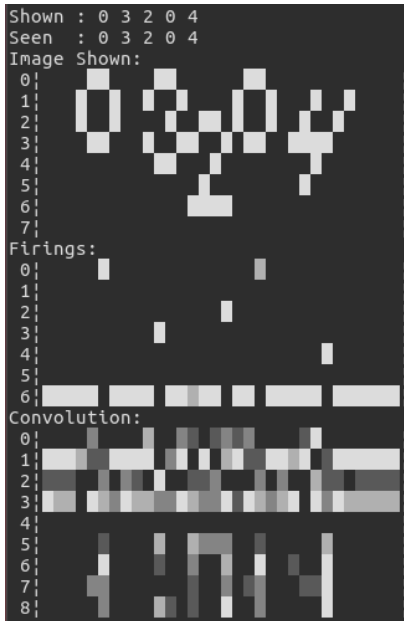


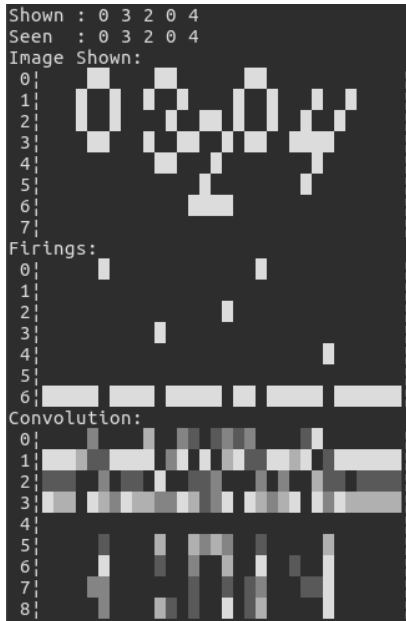


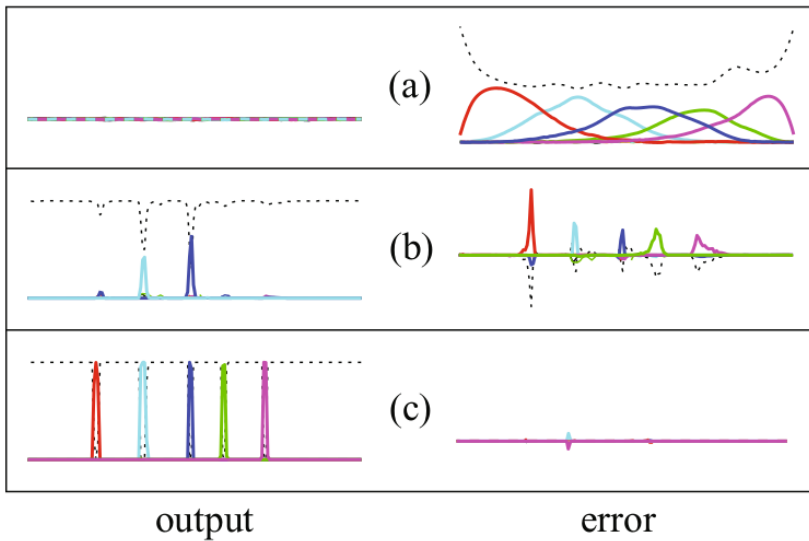












Thank you

- t - h - a - n - k - ' ' - y - o - u -